

Experiment 7: Build a Multiple Linear Regression Model and Predict

the Dependent Variable using the Gradient Descent Method

ANJALI K - 24MCAR0186

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
from sklearn.datasets import fetch_california_housing
```

```
In [8]: # =====
# Step 1: Load and Preprocess Dataset
# =====

# Load California Housing dataset
california = fetch_california_housing()
data = pd.DataFrame(california.data, columns=california.feature_names)
data['PRICE'] = california.target

print("Sample Data:")
print(data.head())

# Handle missing values (if any)
data = data.dropna()

# Select features and target
```

```

X = data[['MedInc', 'AveRooms', 'HouseAge', 'AveOccup']] # selecting a few features for simplicity
y = data['PRICE']

# Normalize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2, random_state=42)

```

Sample Data:

| | MedInc | HouseAge | AveRooms | AveBedrms | Population | AveOccup | Latitude | Longitude | PRICE |
|---|--------|----------|----------|-----------|------------|----------|----------|-----------|-------|
| 0 | 8.3252 | 41.0 | 6.984127 | 1.023810 | 322.0 | 2.555556 | 37.88 | -122.23 | 4.526 |
| 1 | 8.3014 | 21.0 | 6.238137 | 0.971880 | 2401.0 | 2.109842 | 37.86 | -122.22 | 3.585 |
| 2 | 7.2574 | 52.0 | 8.288136 | 1.073446 | 496.0 | 2.802260 | 37.85 | -122.24 | 3.521 |
| 3 | 5.6431 | 52.0 | 5.817352 | 1.073059 | 558.0 | 2.547945 | 37.85 | -122.25 | 3.413 |
| 4 | 3.8462 | 52.0 | 6.281853 | 1.081081 | 565.0 | 2.181467 | 37.85 | -122.25 | 3.422 |

In [9]:

```

# =====
# Step 2: Implement Gradient Descent for MLR
# =====

```

```

def gradient_descent(X, y, learning_rate=0.01, iterations=1000):
    m, n = X.shape
    X_b = np.c_[np.ones((m, 1)), X] # Add intercept term
    theta = np.zeros((n + 1, 1)) # Initialize coefficients
    y = y.values.reshape(-1, 1)
    cost_history = []

    for i in range(iterations):
        y_pred = X_b.dot(theta)
        error = y_pred - y
        gradients = (1/m) * X_b.T.dot(error)
        theta -= learning_rate * gradients

```

```

        cost = (1/(2*m)) * np.sum(error ** 2)
        cost_history.append(cost)

    return theta, cost_history

# Train model
theta, cost_history = gradient_descent(X_train, y_train, learning_rate=0.05, iterations=2000)

```

In [10]: # =====

```

# Step 3: Model Evaluation
# =====

# Predictions
X_test_b = np.c_[np.ones((X_test.shape[0], 1)), X_test]
y_pred = X_test_b.dot(theta)

# Compute metrics
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = np.sqrt(mean_squared_error(y_test, y_pred))

print("\nModel Coefficients:")
print(f"Intercept: {theta[0][0]:.4f}")
for i, col in enumerate(X.columns):
    print(f"{col}: {theta[i+1][0]:.4f}")

print(f"\nR2 Score: {r2:.4f}")
print(f"MAE: {mae:.4f}")
print(f"RMSE: {rmse:.4f}")

```

Model Coefficients:

Intercept: 2.0683

MedInc: 0.8463

AveRooms: -0.0702

HouseAge: 0.2127

AveOccup: -0.0430


R² Score: 0.4983

MAE: 0.6025

RMSE: 0.8108


```
In [11]: # =====
# Step 4: Plots
# =====

# Cost vs Iterations
plt.figure(figsize=(6,4))
plt.plot(range(len(cost_history)), cost_history, color='blue')
plt.title("Cost (Loss) vs Iterations")
plt.xlabel("Iterations")
plt.ylabel("Cost")
plt.grid(True)
plt.show()
```



```
In [12]: # Predicted vs Actual
plt.figure(figsize=(6,4))
plt.scatter(y_test, y_pred, alpha=0.6, color='blue', label='Predictions')
```


```
plt.plot([y_test.min(), y_test.max()],
         [y_test.min(), y_test.max()],
         color='red', linestyle='--', label='Ideal Line (y = x)')
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Predicted vs Actual Values")
plt.legend()
plt.grid(True)
plt.show()
```



In [15]: # Residual Plot

```
residuals = y_test - y_pred.flatten()
plt.figure(figsize=(6,4))
sns.histplot(residuals, kde=True, color='orange')
plt.title("Residuals Distribution")
```

```
plt.xlabel("Residuals")
plt.show()
```



```
In [14]: # Optional: Scatter 2D (using 1 feature for visualization)
plt.figure(figsize=(6,4))
plt.scatter(X_test[:, 0], y_test, color='red', label='Actual')
plt.scatter(X_test[:, 0], y_pred, color='blue', label='Predicted')
plt.xlabel("Median Income (Scaled)")
plt.ylabel("Price")
plt.title("Regression Fit (2D View)")
plt.legend()
plt.show()
```

Regression Fit (2D View)

